

MAXIMUM CARDINALITY MATCHING FOR BIPARTITE GRAPHS

Samuel Ueltschi, Conradin Roffler, Thomas Meier, Isabelle Roesch

Department of Computer Science
ETH Zürich
Zürich, Switzerland

ABSTRACT

Maximum cardinality matching, focus on existing algorithms and optimize the parallel versions in a highly multi-threaded environment. Focus on Pothén-Fan, reason about performance.

1. INTRODUCTION

Motivation. Graph matching has several applications in computer science, for example the marriage problem or computing the block triangular form (BTF) of a sparse matrix [1]. Bipartite graph matching is also a special case of a network flow problem. As data gets bigger, we are interested in the performance of the algorithms that solve these problems.

Related work. Two parallel algorithms that we are testing on Xeon Phi [2] [3]

2. BACKGROUND: ALGORITHMS FOR MAXIMUM MATCHING IN BIPARTITE GRAPHS

Maximum Matching. Given a graph $G = (V, E)$, where V are the vertices and E the edges, a matching M in the graph is a subset of its edges such that no edge in M shares a vertex with another edge in M . A matching M in G is *maximal* if there is no other matching $M' \neq M$ such that $M \subset M'$. A maximal matching M is a *maximum* matching of the graph, if there is no other matching M' such that $|M'| > |M|$.

State of the Art. Best algorithms to solve this at the moment (parallel and sequential), as well as $O(\dots)$ considerations

Initial Matching. Explain greedy matching, enhanced greedy matching (ours) and karp-siper initial matching.

3. ALGORITHMS AND OPTIMIZATIONS

Focus on Pothén-Fan [2] but also report Tree Grafting [3] for completeness

3.1. Pothén-Fan

Parallel Pothén-Fan. Pseudocode for parallel ppf

PRAM Analysis. Show DAG, worst case $O(n)$, best case $O(1)$ with n processors (n nodes), but real world graph are rather $O(1)$

Roofline Model. number of operations, number of moves, what if whole graph fits into cache, etc

Optimizations. Test and Test and Set, Locality, Use only half of the visited array, set only half of the matching vector while setting the rest last, etc

3.2. Tree Grafting

Paper [3] claims that PPF is not well suited for many thin cores.

Parallel Tree Grafting. Explain parallel tree grafting.

Optimizations. Same optimizations as PPF version 3 for the augmenting paths. Non-blocking queue as a data structure.

4. EXPERIMENTAL RESULTS

Experimental setup.

Xeon Phi (5110P), GCC, -O3, 60 simplified Intel CPU cores running at 1056 MHz and supports 4 threads per core, resulting in a total of 240 threads. Each core has a 32kb L1 data cache, a 32kb L1 instruction cache and a private 512 kb L2 unified cache. [4]

Test Data. To test our algorithms, we have used several graphs from real-world examples. The graphs and their attributes are listed in 1. **Benchmarks.**

Sequential Pothén-Fan

Verification.

We use the *Edmonds Maximum Cardinality Matching* algorithm [5] from the Boost Graph library to verify the correctness of our implementations.

Results.

	$ V $	$ E $	Density
coPaperDBLP	1'080'872	15'245'732	5.22e−5
Wikipedia	7'030'396	45'030'392	3.64e−6
Amazon0312	801'454	3'200'440	1.99e−5
Gnutella	73'364	176'656	1.31e−4

Table 1. Test data used for benchmarks. $|V|$ is the number of vertices, $|E|$ the number of edges. The density describes the sparseness of the graph, where 0.0 represents an empty graph and 1.0 a fully connected bipartite graph.

5. CONCLUSIONS

Super linear speedup because of caching effects
PPF scales well on Xeon Phi
Tree Grafting results from the paper could not be reproduced.

6. REFERENCES

- [1] Alex Pothén and Chin-Ju Fan, “Computing the block triangular form of a sparse matrix,” *ACM Trans. Math. Softw.*, vol. 16, no. 4, pp. 303–324, Dec. 1990.
- [2] A. Azad, M. Halappanavar, S. Rajamanickam, E. G. Boman, A. Khan, and A. Pothén, “Multithreaded algorithms for maximum matching in bipartite graphs,” in *2012 IEEE 26th International Parallel and Distributed Processing Symposium*, May 2012, pp. 860–872.
- [3] Ariful Azad, Aydin Bulu, and Alex Pothén, “A parallel tree grafting algorithm for maximum cardinality matching in bipartite graphs,” in *Proceedings of the 2015 IEEE International Parallel and Distributed Processing Symposium*, Washington, DC, USA, 2015, IPDPS ’15, pp. 1075–1084, IEEE Computer Society.
- [4] S. Ramos and T. Hoefer, “Modeling Communication in Cache-Coherent SMP Systems - A Case-Study with Xeon Phi,” in *Proceedings of the 22nd international symposium on High-performance parallel and distributed computing*. Jun. 2013, pp. 97–108, ACM.
- [5] “Boost edmonds: Maximum cardinality matching,” http://www.boost.org/doc/libs/1_36_0/libs/graph/doc/maximum_matching.html, Accessed: 2017-01-05.